



CANDIDATE ENTRY PLAN

First 120 Days

Director of Product Operations, Artificial Intelligence and Machine Learning (AI/ML)

Russell Dudek - candidate-created discovery hypothesis

Operating intent: read the production system and earn trust before standardizing; define a lightweight Model Reliability Contract; prove one reliability path end to end; scale reusable controls without scaling bureaucracy.

ENTRY PRINCIPLES

Observe before prescribing

Baseline before promising

One path before platform-wide rules

Human authority remains explicit

Business outcome joins model health

Exceptions are designed, not hidden

REVIEW 01 / DAYS 1-30

Read the production system.

Objective: understand the actual model lifecycle, ownership seams, platform toil, and consequence landscape.

- Inventory production models, pipelines, critical decisions, owners, data dependencies, release paths, incidents, vendors, and platform debt.
- Shadow data science, data engineering, platform operations, IT/security, product, and business workflow owners.
- Map where release evidence, monitoring, intervention authority, and business-outcome ownership are explicit or missing.
- Baseline deployment lead time, monitoring coverage, alert quality, detection/recovery, cost visibility, and adoption evidence.

Exit evidence: model/service map - ownership map - baseline scorecard - trust-debt register.

REVIEW 02 / DAYS 31-60

Define the contract.

Objective: agree on the smallest useful reliability standard and the conditions that justify exceptions.

- Co-design minimum release evidence by model consequence and use-case class.
- Define a practical service envelope covering quality, drift, data health, latency, availability, freshness, and cost.
- Name intervention rights for tune, retrain, rollback, fallback, pause, and escalation.
- Select one production path where reliability friction and business consequence are visible.

Exit evidence: Contract v0.1 - intervention matrix - prioritized reliability debt - pilot charter.

DISCOVERY BASELINE

Domain	What to trace	Baseline output
Model estate	Critical decisions, consequence, owner, lifecycle stage, dependencies	Model/service map
Release flow	Environments, approvals, manual gates, rework, failure and rollback paths	Lead-time and friction map
Observability	Quality, drift, data health, latency, cost, adoption, business outcome	Coverage and trust-debt register
Ownership	Who may release, tune, retrain, pause, rollback, accept exception, retire	Decision-rights map
Operating economics	Recurring toil, cloud/inference cost, duplicated pipelines, support burden	Prioritized reliability debt

FIRST-MONTH OPERATING RHYTHM

Fieldwork

Shadow one model journey, one incident path, and one business workflow each week.

Friday reliability review

What changed, what surprised us, what needs an owner, and what evidence is missing?

Decision log

Record assumptions, intervention rights, exceptions, and questions that require executive resolution.



CANDIDATE ENTRY PLAN / CONTINUED

Prove, Scale, and Govern

Days 61-120 and the first operating cadence

REVIEW 03 / DAYS 61-90

Prove one path end to end.

Objective: show that the contract improves a real production path without documentation theater.

- Apply the contract with accountable technical and business owners.
- Automate appropriate release checks, monitoring, alert routing, decision evidence, and rollback/fallback preparation.
- Exercise at least one controlled failure or recovery path.
- Measure alert quality, time to detect, time to decide, time to recover, deployment flow, cost, adoption, and business-outcome linkage.
- Capture what should become a platform primitive, a documented exception, or a stopped practice.

Exit evidence: validated path - exercised recovery - measured friction removed - reusable learning record.

REVIEW 04 / DAYS 91-120

Scale the pattern, not the bureaucracy.

Objective: turn proven controls into reusable capability, a credible roadmap, and a shared operating review.

- Convert proven checks and runbooks into platform templates with explicit exception logic.
- Establish a monthly reliability review tied to model health, workflow outcomes, economics, and trust debt.
- Prioritize platform debt by business consequence and recurring toil rather than technical elegance alone.
- Build the capability, hiring, training, vendor, and tooling roadmap from observed constraints.
- Present the next two quarters as choices: standardize, invest, partner, defer, or retire.

Exit evidence: reusable standard - operating review - platform roadmap - talent and capability plan.

PROPOSED OPERATING REVIEW - DISCOVERY HYPOTHESIS

Question	Evidence	Decision
Which services are outside their agreed contract?	Quality, drift, latency, data health, cost, incidents, overrides	Intervene - accept exception - redesign - retire
Where is release flow slow or unreliable?	Lead time, failure rate, rework, manual gates, environment variance	Automate - standardize - simplify - add evidence
Which models are technically healthy but operationally weak?	Adoption, workflow friction, business KPI, human override, trust	Reframe objective - redesign workflow - retrain - stop
What learning should become shared capability?	Recurring incidents, reusable checks, runbooks, toil patterns	Template - platform primitive - training - ownership change

Boundary: model inventory, architecture, priority use cases, metrics, tool choices, and decision rights must be discovered with Southern Glazer's. This plan proposes a method, not a claim about current conditions.

DAY-120 SCORECARD

Dimension	Evidence of progress
Release flow	One path has clearer evidence, less rework, known failure modes, and a tested recovery path.
Reliability	Signal coverage, alert quality, time to decide, and time to recover are measurable.
Ownership	Technical and business intervention rights are explicit and exercised.
Value and adoption	Model health connects to a workflow outcome, user behavior, and accountable owner.
Learning	Incident lessons become templates, platform primitives, training, or stopped practices.