



INTERVIEW THESIS BRIEF

Make AI dependable enough to operate the business.

Russell Dudek / Director of Product Operations, AI/ML candidate thesis

Actual mandate

Turn ML experimentation into a dependable production capability by building the operating system around deployment automation, observability, performance optimization, risk controls, ownership, and cross-functional decision cadence.

Company-specific hypothesis: make AI dependable enough to support widening digital, commercial, and operational decision surfaces - not a collection of models that work only in notebooks.

Company moment

- The role centers on MLOps, monitoring, optimization, tuning, platform operations, cloud, risk, and CI/CD.
- Proof publicly reports 6,000+ daily orders, 32 active markets, and 18,000+ products.
- Southern Glazer's describes advanced automation, connected distribution, tracking, last-mile, and route planning.
- July 2026 reporting describes a hybrid field/inside/digital sales model and AI-assisted adaptation.

Operating tensions

1. Experiment velocity vs production reliability.
2. Shared platform standards vs diverse model and use-case needs.
3. Automation vs human intervention authority.
4. Model metrics vs business outcomes and adoption.
5. Platform sophistication vs operating burden.

White-space thesis

AI earns scale when every model carries an operating contract: what it may influence, how health is measured, who owns intervention, and what evidence justifies the next release.

The contract joins intent and ownership, release evidence, live service envelope, intervention rights, and a learning record.

Proof map

- **Compunetics:** high-reliability systems, quality controls, root cause, technical leadership.
- **Amazon:** 24/7 visibility, escalation, standard work, operating cadence at scale.
- **Vape-Jet:** telemetry-to-action loops across factory, field, customer, support, engineering, and approximately 1,000 fielded robots.
- **DudeWorth:** AI readiness, evaluation, human judgment, governance, adoption, and value.

Strongest objection

The evidence does not establish seven years of ML platform operations or four years leading that exact function. The degree is not presented as one of the specified computing/electronics engineering degrees.

Response: do not pretend equivalence. Position the candidacy as reliability and operating-system leadership around a technically deep MLOps organization.

Second constraint: Pittsburgh location vs required residence near Miramar or Dallas.

Stakeholder topology

Data science; data engineering; platform operations; cloud/IT; security, risk, compliance, and data governance; product/value orchestration; commercial and supply-chain workflow owners; frontline users; executive sponsors.

Change burden

Teams must adopt shared release evidence, explicit ownership, model/service health expectations, incident and retraining playbooks, outcome ownership, human-override logic, feedback capture, and recurring value review.



INTERVIEW BRIEF / CONTINUED

Questions, scorecard, and discovery hypotheses

DECISION ARCHITECTURE

Release? Tune? Retrain? Roll back? Fallback? Escalate? Standardize? Stop? The role should make required evidence and intervention authority explicit at each decision.

120-DAY INTENT

1-30: read and baseline. **31-60:** define Contract v0.1 and intervention rights.

61-90: prove one production path. **91-120:** scale reusable controls, operating review, capability plan, and roadmap.

EXECUTIVE QUESTIONS

1. Which model classes or decisions are already business critical, and where does operational confidence lag experimentation?
2. Where does a model become someone else's problem today?
3. Which failure mode has most shaped the current architecture?
4. What authority should this role have to stop a release or trigger rollback?
5. How are model health, workflow adoption, and business outcomes reconciled?
6. Are you hiring primarily for deep MLOps pedigree, or for a leader who can build the operating system around a technically strong platform organization?

CANDIDATE SCORECARD - HYPOTHESES

Dimension	Candidate measure
Release flow	Lead time, frequency, change failure
Reliability	SLO coverage, detection, recovery, alert quality
Business value	Model-linked outcome and adoption
Economics	Training/inference cost, reuse, duplicated pipeline reduction
Trust & risk	Release evidence, intervention ownership, auditability, exception aging
Learning speed	Incident recurrence; lesson-to-control time

DISCOVERY HYPOTHESES

Likely pressure points to test, not assert: ownership at DS/DE/platform/business handoffs; uneven observability; platform standardization; business-metric linkage; rollback authority; cloud/evaluation economics; adoption burden; and talent/capability mix.

SOURCE LIST

Official Job ID 41852; Southern Glazer's Who We Are; About Proof; Supply Chain Details; July 10, 2026 reporting on hybrid sales-model change and AI adaptation.

FOUR OPERATING REVIEWS

1-30

Map model estate, ownership, release flow, incidents, adoption, cost, and trust debt.

31-60

Define Contract v0.1, consequence tiers, service envelope, and intervention rights.

61-90

Prove one production path; exercise recovery; measure friction and outcome linkage.

91-120

Scale reusable controls, operating review, platform roadmap, and capability plan.

Operating principle: standardize what reduces avoidable ambiguity; preserve deliberate exceptions where decision consequence, model class, market context, or human authority genuinely differ.